

statistical models, hypothesis tests, and data exploration



[statsmodels](#) [about.html#about-statsmodels] is a Python module that provides classes and functions for the estimation of many different statistical models, as well as for conducting statistical tests, and statistical data exploration. An extensive list of result statistics are available for each estimator. The results are tested against existing statistical packages to ensure that they are correct. The package is released under the open source Modified BSD (3-clause) license. The online documentation is hosted at [statsmodels.org](https://www.statsmodels.org/) [https://www.statsmodels.org/].

Introduction

`statsmodels` supports specifying models using R-style formulas and `pandas` DataFrames. Here is a simple example using ordinary least squares:

```
In [1]: import numpy as np

In [2]: import statsmodels.api as sm

In [3]: import statsmodels.formula.api as smf

# Load data
In [4]: dat = sm.datasets.get_rdataset("Guerry", "HistData").data

# Fit regression model (using the natural log of one of the
regressors)
```

```
In [5]: results = smf.ols('Lottery ~ Literacy + np.log(Pop1831)',
data=dat).fit()
```

```
# Inspect the results
```

```
In [6]: print(results.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          Lottery    R-squared:
0.348
Model:                  OLS        Adj. R-squared:
0.333
Method:                 Least Squares    F-statistic:
22.20
Date:                   Fri, 21 Feb 2020    Prob (F-statistic):
1.90e-08
Time:                   13:59:15    Log-Likelihood:
-379.82
No. Observations:      86    AIC:
765.6
Df Residuals:          83    BIC:
773.0
Df Model:               2
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t
[0.025 0.975]				

Intercept	246.4341	35.233	6.995	0.000
176.358 316.510				
Literacy	-0.4889	0.128	-3.832	0.000
-0.743 -0.235				
np.log(Pop1831)	-31.3114	5.977	-5.239	0.000
-43.199 -19.424				

```
=====
Omnibus:                3.713    Durbin-Watson:
2.019
Prob(Omnibus):          0.156    Jarque-Bera (JB):
3.394
Skew:                   -0.487    Prob(JB):
0.183
Kurtosis:               3.003    Cond. No.
```

702.

=====

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

You can also use `numpy` arrays instead of formulas:

```
In [7]: import numpy as np
```

```
In [8]: import statsmodels.api as sm
```

```
# Generate artificial data (2 regressors + constant)
```

```
In [9]: nobs = 100
```

```
In [10]: X = np.random.random((nobs, 2))
```

```
In [11]: X = sm.add_constant(X)
```

```
In [12]: beta = [1, .1, .5]
```

```
In [13]: e = np.random.random(nobs)
```

```
In [14]: y = np.dot(X, beta) + e
```

```
# Fit regression model
```

```
In [15]: results = sm.OLS(y, X).fit()
```

```
# Inspect the results
```

```
In [16]: print(results.summary())
```

OLS Regression Results

=====

Dep. Variable:	y	R-squared:	
0.337			
Model:	OLS	Adj. R-squared:	
0.323			
Method:	Least Squares	F-statistic:	
24.66			
Date:	Fri, 21 Feb 2020	Prob (F-statistic):	
2.19e-09			
Time:	13:59:15	Log-Likelihood:	

```

-12.084
No. Observations:          100    AIC:
30.17
Df Residuals:              97    BIC:
37.98
Df Model:                  2
Covariance Type:          nonrobust
=====

              coef    std err          t      P>|t|      [0.025
0.975]
-----
const         1.4774      0.072     20.651      0.000      1.335
1.619
x1             0.0900      0.098      0.921      0.359     -0.104
0.284
x2             0.6260      0.090      6.960      0.000      0.447
0.804
=====

Omnibus:                26.611    Durbin-Watson:
2.322
Prob(Omnibus):            0.000    Jarque-Bera (JB):
5.863
Skew:                  -0.155    Prob(JB):
0.0533
Kurtosis:              1.855    Cond. No.
5.18
=====

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is
correctly specified.

```

Have a look at `dir(results)` to see available results. Attributes are described in `results.__doc__` and results methods have their own docstrings.

Citation

Please use following citation to cite statsmodels in scientific publications:

Seabold, Skipper, and Josef Perktold. “[statsmodels: Econometric and statistical modeling with python](#).”

[<http://conference.scipy.org/proceedings/scipy2010/pdfs/seabold.pdf>]

Proceedings of the 9th Python in Science Conference. 2010.

Bibtex entry:

```
@inproceedings{seabold2010statsmodels,  
  title={statsmodels: Econometric and statistical modeling with  
python},  
  author={Seabold, Skipper and Perktold, Josef},  
  booktitle={9th Python in Science Conference},  
  year={2010},  
}
```

[Installing statsmodels](#) [install.html]

[Getting started](#) [gettingstarted.html]

[User Guide](#) [user-guide.html]

[Examples](#) [examples/index.html]

[API Reference](#) [api.html]

[About statsmodels](#) [about.html]

[Developer Page](#) [dev/index.html]

[Release Notes](#) [release/index.html]

Index

[Index](#) [genindex.html]

[Module Index](#) [py-modindex.html]